

# Linux Switching – Interconnecting Namespaces

📅 16/09/2013 📁 Linux Networking, Openvswitch

Switching in software on Linux is one of the important parts when using virtualization technologies like KVM or LXC. Typical hosts do not provide one or more physical adapters for each NIC of a virtual machine in KVM or per container when using LXC. Something else must take the part to interconnect the virtual network interfaces.

The software switching classical tool is the linuxbridge, which is available in the Linux kernel for a long time. The frontend to manage the linuxbridge is **brctl**. The newer tool is the Openvswitch (at <http://openvswitch.org/>). The main frontend is **ovs-vsctl**.


In this post I will show multiple solutions to interconnect Linux namespaces using a software based switch. A performance analysis of these solutions will be discussed in another article later. Starting with network namespaces helps to understand the more complex situations when using KVM or LXC.

## tap interfaces

Linux tap interfaces created with `ip tuntap` cannot be used to attach network namespaces to linuxbridges or the openvswitch.

## veth pair

The simple solution to connect two network namespaces is the usage of one veth pair. This has been discussed in a previous article.




Connecting namespaces using a veth pair

The command sequence has been discussed in a previous article, but we show the commands here again

```
veth pair Shell
1 # add the namespaces
2 ip netns add ns1
3 ip netns add ns2
4 # create the veth pair
5 ip link add tap1 type veth peer name tap2
6 # move the interfaces to the namespaces
7 ip link set tap1 netns ns1
8 ip link set tap2 netns ns2
9 # bring up the links
10 ip netns exec ns1 ip link set dev tap1 up
11 ip netns exec ns2 ip link set dev tap2 up
12 # now assign the ip addresses
```

## linux bridge and two veth pairs

When more than two network namespaces (or KVM or LXC instances) must be connected a switch should be used. Linux offers as one solution the well known linux bridge.



Connecting namespaces using a linux bridge and two veth pairs

We need for this setup one switch, and two connectors. In this setup we use a linuxbridge and two veth pairs.

The commands to create this setup are:

```
linuxbridge and two veth pairs Shell
1 # add the namespaces
2 ip netns add ns1
3 ip netns add ns2
4 # create the switch
5 BRIDGE=br-test
```


```

6 brctl addbr $BRIDGE
7 brctl stp $BRIDGE off
8 ip link set dev $BRIDGE up
9 #
10 ##### PORT 1
11 # create a port pair
12 ip link add tap1 type veth peer name br-tap1
13 # attach one side to linuxbridge
14 brctl addif br-test br-tap1
15 # attach the other side to namespace
16 ip link set tap1 netns ns1
17 # set the ports to up
18 ip netns exec ns1 ip link set dev tap1 up
19 ip link set dev br-tap1 up
20 #
21 ##### PORT 2
22 # create a port pair
23 ip link add tap2 type veth peer name br-tap2
24 # attach one side to linuxbridge
25 brctl addif br-test br-tap2
26 # attach the other side to namespace
27 ip link set tap2 netns ns2
28 # set the ports to up
29 ip netns exec ns2 ip link set dev tap2 up
30 ip link set dev br-tap2 up
31 #

```

## openvswitch and two veth pairs

Another solution is to use the openvswitch instead of the „old“ linuxbridge. The configuration is nearly the same as for the linuxbridge.



Connecting namespaces using the openvswitch and two veth pairs

We need for this setup one switch, and two connectors. In this setup we use an openvswitch and


two veth pairs.

The commands to create this setup are:

```
openvswitch and two veth pairs Shell
1 # add the namespaces
2 ip netns add ns1
3 ip netns add ns2
4 # create the switch
5 BRIDGE=ovs-test
6 ovs-vsctl add-br $BRIDGE
7 #
8 ##### PORT 1
9 # create a port pair
10 ip link add tap1 type veth peer name ovs-tap1
11 # attach one side to ovs
12 ovs-vsctl add-port $BRIDGE ovs-tap1
13 # attach the other side to namespace
14 ip link set tap1 netns ns1
15 # set the ports to up
16 ip netns exec ns1 ip link set dev tap1 up
17 ip link set dev ovs-tap1 up
18 #
19 ##### PORT 2
20 # create a port pair
21 ip link add tap2 type veth peer name ovs-tap2
22 # attach one side to ovs
23 ovs-vsctl add-port $BRIDGE ovs-tap2
24 # attach the other side to namespace
25 ip link set tap2 netns ns2
26 # set the ports to up
27 ip netns exec ns2 ip link set dev tap2 up
28 ip link set dev ovs-tap2 up
29 #
```

## openvswitch and two openvswitch ports

Another solution is to use the openvswitch and make use of the openvswitch internal ports. This avoids the usage of the veth pairs, which must be used in all other solutions.



Connecting namespaces using the openvswitch and two openvswitch ports

We need for this setup one switch, and two connectors. In this setup we use an openvswitch and two openvswitch ports.

The commands to create this setup are:

openvswitch and two openvswitch internal ports

Shell

```

1 # add the namespaces
2 ip netns add ns1
3 ip netns add ns2
4 # create the switch
5 BRIDGE=ovs-test
6 ovs-vsctl add-br $BRIDGE
7 #
8 ##### PORT 1
9 # create an internal ovs port
10 ovs-vsctl add-port $BRIDGE tap1 -- set Interface tap1 type=internal
11 # attach it to namespace
12 ip link set tap1 netns ns1
13 # set the ports to up
14 ip netns exec ns1 ip link set dev tap1 up
15 #
16 ##### PORT 2
17 # create an internal ovs port
18 ovs-vsctl add-port $BRIDGE tap2 -- set Interface tap2 type=internal
19 # attach it to namespace
20 ip link set tap2 netns ns2
21 # set the ports to up
22 ip netns exec ns2 ip link set dev tap2 up

```

## Performance

In another article I will show some performance numbers for the four presented solutions. There are noticeable differences with respect to throughput and CPU usage.

